

Bike Store Sales

October 30, 2025

Bike Store Sales

Hands on!

```
[4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline
```

Loading the data:

```
[6]: sales = pd.read_csv(r"C:\Users\Admin\Downloads\sales_data.csv",
    ↪ parse_dates=["Date"])
```

The data glance:

```
[7]: sales.head()
```

```
[7]:
```

	Date	Day	Month	Year	Customer_Age	Age_Group	\
0	2013-11-26	26	November	2013	19	Youth (<25)	
1	2015-11-26	26	November	2015	19	Youth (<25)	
2	2014-03-23	23	March	2014	49	Adults (35-64)	
3	2016-03-23	23	March	2016	49	Adults (35-64)	
4	2014-05-15	15	May	2014	47	Adults (35-64)	

	Customer_Gender	Country	State	Product_Category	Sub_Category	\
0	M	Canada	British Columbia	Accessories	Bike Racks	
1	M	Canada	British Columbia	Accessories	Bike Racks	
2	M	Australia	New South Wales	Accessories	Bike Racks	
3	M	Australia	New South Wales	Accessories	Bike Racks	
4	F	Australia	New South Wales	Accessories	Bike Racks	

	Product	Order_Quantity	Unit_Cost	Unit_Price	Profit	Cost	\
0	Hitch Rack - 4-Bike	8	45	120	590	360	
1	Hitch Rack - 4-Bike	8	45	120	590	360	
2	Hitch Rack - 4-Bike	23	45	120	1366	1035	
3	Hitch Rack - 4-Bike	20	45	120	1188	900	
4	Hitch Rack - 4-Bike	4	45	120	238	180	

	Revenue
0	950
1	950
2	2401
3	2088
4	418

```
[8]: sales.shape
```

```
[8]: (113036, 18)
```

```
[9]: sales.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 113036 entries, 0 to 113035
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  113036 non-null  datetime64[ns]
1   Day                   113036 non-null  int64
2   Month                 113036 non-null  object
3   Year                  113036 non-null  int64
4   Customer_Age          113036 non-null  int64
5   Age_Group             113036 non-null  object
6   Customer_Gender       113036 non-null  object
7   Country               113036 non-null  object
8   State                 113036 non-null  object
9   Product_Category      113036 non-null  object
10  Sub_Category          113036 non-null  object
11  Product               113036 non-null  object
12  Order_Quantity        113036 non-null  int64
13  Unit_Cost             113036 non-null  int64
14  Unit_Price            113036 non-null  int64
15  Profit               113036 non-null  int64
16  Cost                 113036 non-null  int64
17  Revenue              113036 non-null  int64
dtypes: datetime64[ns](1), int64(9), object(8)
memory usage: 15.5+ MB
```

```
[10]: sales.describe()
```

```
[10]:
```

	Date	Day	Year \
count	113036	113036.000000	113036.000000
mean	2014-11-23 12:14:55.063519232	15.665753	2014.401739
min	2011-01-01 00:00:00	1.000000	2011.000000
25%	2013-12-22 00:00:00	8.000000	2013.000000
50%	2014-06-27 00:00:00	16.000000	2014.000000
75%	2016-01-09 00:00:00	23.000000	2016.000000

max	2016-07-31 00:00:00	31.000000	2016.000000
std	NaN	8.781567	1.272510

	Customer_Age	Order_Quantity	Unit_Cost	Unit_Price \
count	113036.000000	113036.000000	113036.000000	113036.000000
mean	35.919212	11.901660	267.296366	452.938427
min	17.000000	1.000000	1.000000	2.000000
25%	28.000000	2.000000	2.000000	5.000000
50%	35.000000	10.000000	9.000000	24.000000
75%	43.000000	20.000000	42.000000	70.000000
max	87.000000	32.000000	2171.000000	3578.000000
std	11.021936	9.561857	549.835483	922.071219

	Profit	Cost	Revenue
count	113036.000000	113036.000000	113036.000000
mean	285.051665	469.318695	754.370360
min	-30.000000	1.000000	2.000000
25%	29.000000	28.000000	63.000000
50%	101.000000	108.000000	223.000000
75%	358.000000	432.000000	800.000000
max	15096.000000	42978.000000	58074.000000
std	453.887443	884.866118	1309.094674

Numerical analysis and visualization

We'll analyze the Unit_Cost column:

```
[11]: sales['Unit_Cost'].describe()
```

```
[11]: count    113036.000000
      mean      267.296366
      std       549.835483
      min        1.000000
      25%        2.000000
      50%        9.000000
      75%       42.000000
      max      2171.000000
      Name: Unit_Cost, dtype: float64
```

```
[12]: sales['Unit_Cost'].mean()
```


```
[12]: np.float64(267.296365759581)
```

```
[13]: sales['Unit_Cost'].median()
```

```
[13]: 9.0
```

```
[14]: sales['Unit_Cost'].plot(kind='box', vert=False, figsize=(14,6))
plt.show()
```


```
# What I can see ?
# The Unit_Cost distribution is right-skewed - most products are low-cost,
# with a few high-cost outliers. The median represents the typical cost better
  ↳ than the mean.
# Outliers may indicate premium products or data issues.
# A log transformation could help normalize the data for further analysis.
```



```
[15]: sales['Log_Unit_Cost']=np.log(sales['Unit_Cost']+1)
```


```
[16]: sales['Log_Unit_Cost'].plot(kind='box' , vert = False , figsize =(14,6))
plt.show()
```

```
# Applied log transformation to reduce right skewness in Unit_Cost.
# This helps compress large values, making the distribution more symmetrical.
# It improves visibility of lower-cost products and prepares the data for
  ↳ accurate statistical analysis
```



```
[17]: sales['Unit_Cost'].plot(kind = 'density' , figsize =(14,6)) # kde (Kernel
      ↪ Density Estimation plot)
      plt.show()
```

*# The KDE curve shows the distribution of Unit_Cost values.
 # Most products have low costs, with a long right tail indicating a few
 ↪ expensive items.*




```
[18]: sales[sales['Unit_Cost'] < 0]
      # The curve near -1000 is just a smoothing side effect, not an actual data
      ↪ problem.
```

```
[18]: Empty DataFrame
      Columns: [Date, Day, Month, Year, Customer_Age, Age_Group, Customer_Gender,
```

Country, State, Product_Category, Sub_Category, Product, Order_Quantity,
Unit_Cost, Unit_Price, Profit, Cost, Revenue, Log_Unit_Cost]
Index: []

```
[19]: ax = sales['Unit_Cost'].plot(kind='density', figsize=(14,6))  
ax.set_xlim(0, sales['Unit_Cost'].max()) # start from 0 / remove negative  
      ↪ areas  
plt.show()
```




```
[20]: ax = sales['Unit_Cost'].plot(kind='density', figsize=(14,6))  
ax.axvline(sales['Unit_Cost'].mean(),color='red',linestyle='--',label='mean')  
ax.axvline(sales['Unit_Cost'].median(),color='green',linestyle='--',label =  
      ↪ 'median')  
ax.legend()  
plt.show()
```

*#The KDE chart shows a right-skewed distribution, where most unit costs are low,
↪ with a few high-cost outliers.*


#The mean (red line) is slightly higher than the median (green line).

*#This indicates that the average cost is influenced by a small number of
↪ expensive products.*



```
[21]: ax = sales['Unit_Cost'].plot(kind='hist', figsize=(14,6))
ax.set_ylabel('Number Of Items')
ax.set_xlabel('Dollars $')
plt.show()
```

*#Tall bars on the left => many cheap items.
#Long tail on the right => few expensive items*



Categorical analysis and visualization

We'll analyze the Age_Group column:

```
[22]: sales.head()
```

```
[22]:
```

	Date	Day	Month	Year	Customer_Age	Age_Group	\
0	2013-11-26	26	November	2013	19	Youth (<25)	
1	2015-11-26	26	November	2015	19	Youth (<25)	
2	2014-03-23	23	March	2014	49	Adults (35-64)	
3	2016-03-23	23	March	2016	49	Adults (35-64)	
4	2014-05-15	15	May	2014	47	Adults (35-64)	

	Customer_Gender	Country	State	Product_Category	Sub_Category	\
0	M	Canada	British Columbia	Accessories	Bike Racks	
1	M	Canada	British Columbia	Accessories	Bike Racks	
2	M	Australia	New South Wales	Accessories	Bike Racks	
3	M	Australia	New South Wales	Accessories	Bike Racks	
4	F	Australia	New South Wales	Accessories	Bike Racks	

	Product	Order_Quantity	Unit_Cost	Unit_Price	Profit	Cost	\
0	Hitch Rack - 4-Bike	8	45	120	590	360	
1	Hitch Rack - 4-Bike	8	45	120	590	360	
2	Hitch Rack - 4-Bike	23	45	120	1366	1035	
3	Hitch Rack - 4-Bike	20	45	120	1188	900	
4	Hitch Rack - 4-Bike	4	45	120	238	180	


	Revenue	Log_Unit_Cost
0	950	3.828641
1	950	3.828641
2	2401	3.828641
3	2088	3.828641
4	418	3.828641

```
[23]: sales['Age_Group'].value_counts()
```

```
[23]: Age_Group
Adults (35-64)          55824
Young Adults (25-34)    38654
Youth (<25)             17828
Seniors (64+)           730
Name: count, dtype: int64
```


```
[24]: sales['Age_Group'].value_counts().plot(kind='pie',figsize=(6,6))
plt.show()

# Pie Chart Displays the percentage distribution of Age_Group
```

```
[25]: ax = sales['Age_Group'].value_counts().plot(kind='bar',figsize=(14,6))
      ax.set_ylabel('Number Of Sales')
      plt.show()

      # Bar Chart identify which customer age group contributes most to total sales
```



Relationship between the columns?

Can we find any significant relationship?

```
[26]: sales.drop(columns=['Log_Unit_Cost'] , inplace=True)
```

```
[27]: correlation = sales.select_dtypes(include=['number']).corr()
correlation

# What I Found ?
# - Unit_Cost and Unit_Price are high positive correlated (~0.99), meaning
  ↳ prices rise with cost.
# - Profit shows a strong positive correlation with both Revenue (~0.96) and
  ↳ Cost (~0.90),
# means that higher sales increase both costs and profits
# - Order_Quantity has a moderate negative correlation with Unit_Cost and
  ↳ Unit_Price (~-0.5),
# suggesting bulk purchases come with lower unit prices.
# - Customer_Age, Day, and Year show weak correlations with other variables,
# indicating little linear relationship with sales metrics.
```

```
[27]:
```

	Day	Year	Customer_Age	Order_Quantity	Unit_Cost	\
Day	1.000000	-0.007635	-0.014296	-0.002412	0.003133	
Year	-0.007635	1.000000	0.040994	0.123169	-0.217575	
Customer_Age	-0.014296	0.040994	1.000000	0.026887	-0.021374	
Order_Quantity	-0.002412	0.123169	0.026887	1.000000	-0.515835	
Unit_Cost	0.003133	-0.217575	-0.021374	-0.515835	1.000000	

Unit_Price	0.003207	-0.213673	-0.020262	-0.515925	0.997894
Profit	0.004623	-0.181525	0.004319	-0.238863	0.741020
Cost	0.003329	-0.215604	-0.016013	-0.340382	0.829869
Revenue	0.003853	-0.208673	-0.009326	-0.312895	0.817865

	Unit_Price	Profit	Cost	Revenue
Day	0.003207	0.004623	0.003329	0.003853
Year	-0.213673	-0.181525	-0.215604	-0.208673
Customer_Age	-0.020262	0.004319	-0.016013	-0.009326
Order_Quantity	-0.515925	-0.238863	-0.340382	-0.312895
Unit_Cost	0.997894	0.741020	0.829869	0.817865
Unit_Price	1.000000	0.749870	0.826301	0.818522
Profit	0.749870	1.000000	0.902233	0.956572
Cost	0.826301	0.902233	1.000000	0.988758
Revenue	0.818522	0.956572	0.988758	1.000000

```
[28]: fig = plt.figure(figsize=(8,8))
plt.matshow(corr, cmap='RdBu', fignum=0)
plt.xticks(range(len(corr.columns)), corr.columns, rotation='vertical');
plt.yticks(range(len(corr.columns)), corr.columns);
plt.show()

# In this heatmap
# Blue => Positive correlation (variables increase together)
# Red => Negative correlation (one increases while the other decreases)
# Light color => Weak or no correlation (little relationship)

# The heatmap shows strong positive correlations among Profit, Revenue, and
# Cost -
# meaning profitability rises with higher sales and expenses.
# Order_Quantity is negatively correlated with Unit_Price/Cost, suggesting bulk
# discounts.
# Time and Customer_Age have little direct impact on financial metrics.
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[28], line 2
      1 fig = plt.figure(figsize=(8,8))
----> 2 plt.matshow(corr, cmap='RdBu', fignum=0)
      3 plt.xticks(range(len(corr.columns)), corr.columns, rotation='vertical')
      4 plt.yticks(range(len(corr.columns)), corr.columns);

NameError: name 'corr' is not defined
```

```
[ ]: sales.plot(kind='scatter' , x='Customer_Age' , y='Revenue' , figsize=(6,6))
plt.show()
```

```
# The scatter plot shows that most revenue comes from customers aged 20-50.
# Older customers (60+) generate less revenue, and there's no clear correlation
# between age and spending. A few high-value outliers exist around age 50.
```

```
[ ]: sales.plot(kind='scatter' , x='Revenue' , y='Profit' , figsize=(6,6))
plt.show()

# The scatter plot shows a clear positive relationship between Revenue and
↪ Profit -
# higher sales generally lead to higher profits. A few outliers may indicate
# exceptional or unprofitable transactions.
```

```
[ ]: ax = sales[['Profit', 'Age_Group']].boxplot(by='Age_Group',figsize=(10,6))
ax.set_ylabel('Profit')
plt.show()

# This box plot compares Profit distribution across different Age_Groups.
# It helps identify which age segment generates the most profit
# and highlights variations or outliers within each group.
```

```
[ ]: boxplot_cols = ['Year', 'Customer_Age', 'Order_Quantity', 'Unit_Cost',
↪ 'Unit_Price', 'Profit']
sales[boxplot_cols].plot(kind='box' , subplots=True , layout=(2,3) ,
↪ figsize=(14,8))
plt.show()
```

```
[ ]: sales['Customer_Age'].value_counts()
```

```
[ ]: sales[sales['Profit']<0]
```

Column wrangling

We can also create new columns or modify existing ones.

Add and calculate a new Revenue_per_Age column by dividing Revenue by Customer_Age for each entry in the dataset.

```
[ ]: sales['Revenue_Per_Age'] = sales['Revenue'] / sales['Customer_Age']
sales['Revenue_Per_Age'].head()
```

```
[ ]: sales['Revenue_Per_Age'].plot(kind='density' , figsize=(14,6))
plt.show()

# The sharp peak indicates that most customers contribute a similar amount of
↪ revenue relative to their age,
# suggesting a dominant customer segment or consistent pricing model.
```

```
# The smooth curve helps identify the central tendency and spread, while any  
↳skew or tail could point to outliers-  
# such as older customers with unusually high revenue impact.
```

Add and calculate a new `Calculated_Cost` column

Use This Formula

`Calculated_Cost = Order_Quantity * Unit_Cost`

```
[ ]: sales['Calculated_Cost'] = sales['Order_Quantity'] * sales['Unit_Cost']  
sales['Calculated_Cost'].head()
```

```
[ ]: sales.head(30)
```

```
[ ]: (sales['Calculated_Cost'] != sales['Cost']).sum()
```

Add and calculate a new `Calculated_Revenue` column

Use This Formula

`Calculated_Revenue = Profit + Cost`

```
[ ]: sales['Calculated_Revenue'] = sales['Profit'] + sales['Cost']  
sales['Calculated_Revenue'].head()
```

```
[ ]: (sales['Calculated_Revenue'] != sales['Revenue']).sum()
```

```
[ ]: sales.head()
```

```
[ ]: sales['Revenue'].plot( kind='hist' , bins=100 , figsize=(14,6))  
plt.show()
```

```
[ ]: # Define thresholds using quantiles  
low_threshold = sales['Revenue'].quantile(0.33)  
high_threshold = sales['Revenue'].quantile(0.66)  
  
# Create a segmentation function  
def segment_customer(revenue):  
    if revenue <= low_threshold:  
        return 'Low Revenue'  
    elif revenue <= high_threshold:  
        return 'Medium Revenue'  
    else:  
        return 'High Revenue'  
  
# Apply segmentation  
sales['Revenue_Segment'] = sales['Revenue'].apply(segment_customer)  
  
# Preview the result
```

```
sales[['Revenue', 'Revenue_Segment']].head()
```

```
[ ]: sales.head()
```

```
[ ]: # Count customers in each segment
segment_counts = sales['Revenue_Segment'].value_counts().sort_index()

# Plot the bar chart
plt.figure(figsize=(8, 5))
plt.bar(segment_counts.index, segment_counts.values, color=['#e74c3c', '
↳'#f1c40f', '#2ecc71'])
plt.title('Customer Distribution by Revenue Segment')
plt.xlabel('Revenue Segment')
plt.ylabel('Number of Customers')
plt.tight_layout()
plt.show()
```

Modify all Unit_Price values adding 3% tax to them

new price=original price+(original price×3%)

```
[ ]: sales['Unit_Price_Tax'] = sales['Unit_Price'] * 1.03
sales['Unit_Price_Tax'].head()
```

```
[ ]: sales.head()
```

Selection & Indexing:

Get all the sales made in the state of Kentucky

```
[ ]: sales.loc[sales['State'] == 'Kentucky']
```

```
[ ]: sales.loc[sales['State'] == 'Kentucky'].shape
```

```
[ ]: sales.loc[sales['State'] == 'Kentucky', 'Customer_Gender'].value_counts()
```

Get the mean revenue of the sales group Adults (35-64)

```
[ ]: sales.loc[sales['Age_Group'] == 'Adults (35-64)', 'Revenue'].mean().round(2)
```

How many records belong to Age Group Youth (<25) or Adults (35-64) ?

```
[ ]: sales.loc[(sales['Age_Group'] == 'Youth (<25)') | (sales['Age_Group'] == '
↳'Adults (35-64)')].shape[0]
```

Get the mean revenue of the sales group Adults (35-64) in United States ?

```
[ ]: sales.loc[(sales['Age_Group'] == 'Adults (35-64)') & (sales['Country'] == '
↳'United States'), 'Revenue'].mean().round(2)
```

Increase the revenue by 10% to every sale made in France

```
[ ]: sales.loc[sales['Country'] == 'France', 'Revenue'].head()

[ ]: sales.loc[sales['Country'] == 'France', 'Revenue'] *=1.1

# Increasing France's revenue by 10% to simulate a pricing or growth strategy.
# This helps analyze potential profit impact or forecast future performance.

[ ]: sales['Revenue'] = sales['Revenue'].astype(float)

[ ]: sales.loc[sales['Country'] == 'France', 'Revenue'].head()

[29]: sales['Product'].unique()

[29]: array(['Hitch Rack - 4-Bike', 'All-Purpose Bike Stand',
        'Mountain Bottle Cage', 'Water Bottle - 30 oz.',
        'Road Bottle Cage', 'AWC Logo Cap', 'Bike Wash - Dissolver',
        'Fender Set - Mountain', 'Half-Finger Gloves, L',
        'Half-Finger Gloves, M', 'Half-Finger Gloves, S',
        'Sport-100 Helmet, Black', 'Sport-100 Helmet, Red',
        'Sport-100 Helmet, Blue', 'Hydration Pack - 70 oz.',
        'Short-Sleeve Classic Jersey, XL',
        'Short-Sleeve Classic Jersey, L', 'Short-Sleeve Classic Jersey, M',
        'Short-Sleeve Classic Jersey, S', 'Long-Sleeve Logo Jersey, M',
        'Long-Sleeve Logo Jersey, XL', 'Long-Sleeve Logo Jersey, L',
        'Long-Sleeve Logo Jersey, S', 'Mountain-100 Silver, 38',
        'Mountain-100 Silver, 44', 'Mountain-100 Black, 48',
        'Mountain-100 Silver, 48', 'Mountain-100 Black, 38',
        'Mountain-200 Silver, 38', 'Mountain-100 Black, 44',
        'Mountain-100 Silver, 42', 'Mountain-200 Black, 46',
        'Mountain-200 Silver, 42', 'Mountain-200 Silver, 46',
        'Mountain-200 Black, 38', 'Mountain-100 Black, 42',
        'Mountain-200 Black, 42', 'Mountain-400-W Silver, 46',
        'Mountain-500 Silver, 40', 'Mountain-500 Silver, 44',
        'Mountain-500 Black, 48', 'Mountain-500 Black, 40',
        'Mountain-400-W Silver, 42', 'Mountain-500 Silver, 52',
        'Mountain-500 Black, 52', 'Mountain-500 Silver, 42',
        'Mountain-500 Black, 44', 'Mountain-500 Silver, 48',
        'Mountain-400-W Silver, 38', 'Mountain-400-W Silver, 40',
        'Mountain-500 Black, 42', 'Road-150 Red, 48', 'Road-150 Red, 62',
        'Road-750 Black, 48', 'Road-750 Black, 58', 'Road-750 Black, 52',
        'Road-150 Red, 52', 'Road-150 Red, 44', 'Road-150 Red, 56',
        'Road-750 Black, 44', 'Road-350-W Yellow, 40',
        'Road-350-W Yellow, 42', 'Road-250 Black, 44',
        'Road-250 Black, 48', 'Road-350-W Yellow, 48',
        'Road-550-W Yellow, 44', 'Road-550-W Yellow, 38',
        'Road-250 Black, 52', 'Road-550-W Yellow, 48', 'Road-250 Red, 58',
        'Road-250 Black, 58', 'Road-250 Red, 52', 'Road-250 Red, 48',
```

```

'Road-250 Red, 44', 'Road-550-W Yellow, 42',
'Road-550-W Yellow, 40', 'Road-650 Red, 48', 'Road-650 Red, 60',
'Road-650 Black, 48', 'Road-350-W Yellow, 44', 'Road-650 Red, 52',
'Road-650 Black, 44', 'Road-650 Red, 62', 'Road-650 Red, 58',
'Road-650 Black, 60', 'Road-650 Black, 58', 'Road-650 Black, 52',
'Road-650 Black, 62', 'Road-650 Red, 44',
'Women's Mountain Shorts, M', 'Women's Mountain Shorts, S',
'Women's Mountain Shorts, L', 'Racing Socks, L', 'Racing Socks, M',
'Mountain Tire Tube', 'Touring Tire Tube', 'Patch Kit/8 Patches',
'HL Mountain Tire', 'LL Mountain Tire', 'Road Tire Tube',
'LL Road Tire', 'Touring Tire', 'ML Mountain Tire', 'HL Road Tire',
'ML Road Tire', 'Touring-1000 Yellow, 50', 'Touring-1000 Blue, 46',
'Touring-1000 Yellow, 60', 'Touring-1000 Blue, 50',
'Touring-3000 Yellow, 50', 'Touring-3000 Blue, 54',
'Touring-3000 Blue, 58', 'Touring-3000 Yellow, 44',
'Touring-3000 Yellow, 54', 'Touring-3000 Blue, 62',
'Touring-3000 Blue, 44', 'Touring-1000 Blue, 54',
'Touring-1000 Yellow, 46', 'Touring-1000 Blue, 60',
'Touring-3000 Yellow, 62', 'Touring-1000 Yellow, 54',
'Touring-2000 Blue, 54', 'Touring-3000 Blue, 50',
'Touring-3000 Yellow, 58', 'Touring-2000 Blue, 46',
'Touring-2000 Blue, 50', 'Touring-2000 Blue, 60',
'Classic Vest, L', 'Classic Vest, M', 'Classic Vest, S'],
dtype=object)


```

What's the mean of Customers_Age ?


```
[38]: sales['Customer_Age'].mean().round(0)
```

```
[38]: np.float64(36.0)
```

```
[42]: sales['Customer_Age'].plot(kind='box',vert=False,figsize=(14,6))
plt.show()
```




```
[45]: sales['Customer_Age'].plot(kind='kde',figsize=(14,6))  
plt.show()
```



```
[46]: sales['Year'].value_counts()
```

```
[46]: Year  
2014    29398  
2016    29398  
2013    24443  
2015    24443  
2012     2677  
2011     2677  
Name: count, dtype: int64
```

Can you see any relationship between Unit_Cost and Unit_Price ?

```
[48]: sales['Month'].value_counts()
```

```
[48]: Month  
June      11234  
December  11200  
May       11128  
April     10182  
March      9674  
January    9284  
February   9022  
October    8750  
November   8734  
August     8200
```


```
September    8166
July          7462
Name: count, dtype: int64
```

Can you see any relationship between Order_Quantity and Profit?

```
[63]: sales[['Order_Quantity', 'Profit']].corr()
      # Negative Correlation
```

```
[63]:          Order_Quantity    Profit
Order_Quantity    1.000000 -0.238863
Profit            -0.238863  1.000000
```

```
[64]: sales.sample(100,random_state=1).plot(kind='scatter', x='Order_Quantity',
      ↪ y='Profit',alpha=0.8,figsize=(14,6))
plt.show()
```




Can you see any relationship between Profit and Country?

```
[73]: sales.groupby('Country')['Profit'].sum()
```

```
[73]: Country
Australia    6776030
Canada       3717296
France       2880282
Germany      3359995
United Kingdom 4413853
United States 11073644
Name: Profit, dtype: int64
```

```
[72]: sales[['Profit', 'Country']].boxplot(by='Country', figsize=(10,6))
plt.show()
```



```
[76]: sales['Calculated_Date'] = sales[['Year', 'Month', 'Day']].apply(lambda x :
    ↪ '{}-{}-{}'.format(x[0],x[1],x[2]),axis=1)
sales['Calculated_Date'] = pd.to_datetime(sales['Calculated_Date'])
sales['Calculated_Date'].head()
```


C:\Users\Admin\AppData\Local\Temp\ipykernel_7472\2011521891.py:1: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`

```
sales['Calculated_Date'] = sales[['Year', 'Month', 'Day']].apply(lambda x :
'{}-{}-{}'.format(x[0],x[1],x[2]),axis=1)
```

```
[76]: 0    2013-11-26
      1    2015-11-26
      2    2014-03-23
      3    2016-03-23
      4    2014-05-15
      Name: Calculated_Date, dtype: datetime64[ns]
```

```
[93]: sales.head(100).loc[sales['Year']== 2014].
    ↪ plot(kind='line',x='Month',y='Revenue',figsize=(14,6))
```

```
plt.show()
```



How many Bike Racks orders were made from Canada?

```
[88]: sales.loc[(sales['Country'] == 'Canada') | (sales['Sub_Category'] == 'Bike_Rack')].shape[0]
```

```
[88]: 14178
```


0.0.1 How many orders were made in each region (state) of France?

```
[95]: france_states = sales.loc[sales['Country'] == 'France', 'State'].value_counts()
france_states
```

```
[95]: State
Seine (Paris)      2328
Seine Saint Denis  1684
Nord               1670
Hauts de Seine     1084
Essonne            994
Yveline            954
Seine et Marne     394
Moselle            386
Loiret             382
Val d'Oise         264
Garonne (Haute)    208
Val de Marne       158
Charente-Maritime  148
Somme              134
Loir et Cher       120
```

```
Pas de Calais          90
Name: count, dtype: int64
```

```
[102]: france_states.plot(kind='bar', figsize=(14,6),color='darkblue')
plt.title('Sales Count Per State Of Francr')
plt.xlabel('States')
plt.ylabel('Number Of Sales')
plt.show()
```




0.1 How many orders were made per accessory sub-categories?

```
[104]: accessories = sales.loc[sales['Product_Category']=='Accessories', 'Sub_Category'].value_counts()
accessories
```

```
[104]: Sub_Category
Tires and Tubes      33870
Bottles and Cages    15876
Helmets              12158
Fenders              4032
Cleaners             1802
Hydration Packs      1334
Bike Racks           592
Bike Stands          456
Name: count, dtype: int64
```

```
[105]: accessories.plot(kind='bar', figsize=(14,6),color='darkblue')
plt.show()
```




0.2 How many orders were made per bike sub-categories?

```
[117]: bikes = sales.loc[sales['Product_Category'] == 'Bikes', 'Sub_Category'].
        value_counts()
bikes
```

```
[117]: Sub_Category
Road Bikes      13430
Mountain Bikes   8854
Touring Bikes    3698
Name: count, dtype: int64
```

```
[122]: bikes.plot(kind='pie',figsize=(8,4))
plt.show()
```




```
[130]: def Gender(Gen):  
        if Gen == 'M':  
            return 'Male'  
        else:  
            return 'Female'
```

```
[135]: sales['Customer_Gender'].apply(Gender).value_counts()
```

```
[135]: Customer_Gender  
Male      58312  
Female    54724  
Name: count, dtype: int64
```

```
[132]: sales['Customer_Gender'].apply(Gender).value_counts().  
        plot(kind='bar',figsize=(12,5))  
plt.show()
```



0.3 How many sales with more than 500 in Revenue were made by men?

```
[140]: sales.loc[(sales['Customer_Gender']=='M') & (sales['Revenue'] > 500) ].shape[0]
```

```
[140]: 21773
```

0.4 Get the top-5 sales with the highest revenue

```
[147]: sales.sort_values(['Revenue'], ascending= False ) .head()
```

```
[147]:
```

	Date	Day	Month	Year	Customer_Age	Age_Group	\
112073	2015-07-24	24	July	2015	52	Adults (35-64)	
112072	2013-07-24	24	July	2013	52	Adults (35-64)	
71129	2011-07-08	8	July	2011	22	Youth (<25)	
70307	2011-04-30	30	April	2011	44	Adults (35-64)	
70601	2011-09-30	30	September	2011	19	Youth (<25)	

	Customer_Gender	Country	State	Product_Category	\
112073	M	Australia	Queensland	Clothing	
112072	M	Australia	Queensland	Clothing	
71129	M	Canada	Alberta	Bikes	
70307	M	Canada	British Columbia	Bikes	
70601	F	Canada	British Columbia	Bikes	

	Sub_Category	Product	Order_Quantity	Unit_Cost	\
112073	Vests	Touring-1000 Yellow, 50	29	1482	
112072	Vests	Touring-1000 Yellow, 50	27	1482	
71129	Road Bikes	Road-150 Red, 48	4	2171	
70307	Road Bikes	Road-150 Red, 62	4	2171	

70601	Road Bikes		Road-150 Red, 62		4	2171
-------	------------	--	------------------	--	---	------

	Unit_Price	Profit	Cost	Revenue	Calculated_Date
112073	2384	15096	42978	58074	2015-07-24
112072	2384	14055	40014	54069	2013-07-24
71129	3578	5628	8684	14312	2011-07-08
70307	3578	5485	8684	14169	2011-04-30
70601	3578	5485	8684	14169	2011-09-30

0.5 What is the mean Order_Quantity of orders with more than 10K in revenue?

```
[152]: sales.loc[sales['Revenue']> 10000 , 'Order_Quantity'].mean()
```

```
[152]: np.float64(3.7218934911242605)
```

0.6 How many orders were made in May of 2016?

```
[155]: sales.loc[(sales['Year'] == 2016) & (sales['Month'] == 'May')].shape[0]
```

```
[155]: 5015
```

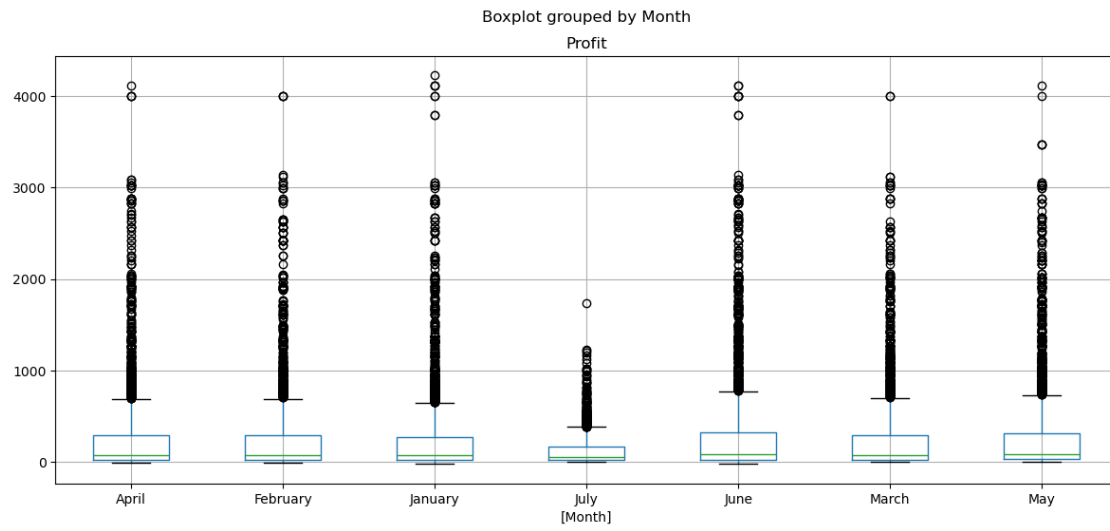
0.7 How many orders were made between May and July of 2016?

```
[158]: cond = (sales['Year']== 2016 ) & (sales['Month'].isin(['May','June','July']))
sales.loc[cond].shape[0]
```

```
[158]: 12164
```

```
[162]: profit_2016 = sales.loc[sales['Year'] == 2016, ['Profit', 'Month']]

profit_2016.boxplot(by='Month', figsize=(14,6))
plt.show()
```



[]: